

Release Notes for Brisk 1.0 beta 1

Introduction to Brisk:

Brisk is an open-source Hadoop and Hive distribution developed by DataStax that utilizes Apache Cassandra for its core services and storage. Brisk provides Hadoop MapReduce capabilities using CassandraFS, an HDFS-compatible storage layer inside Cassandra. By replacing HDFS with CassandraFS, users are able to leverage their current MapReduce jobs on Cassandra's peer-to-peer, fault-tolerant, and scalable architecture. Brisk is also able to support dual workloads, allowing you to use the same cluster of machines for both real-time applications and data analytics without having to move the data around between systems.

Brisk is now available via Apache license v2.0. The DataStax team welcomes your valued feedback.

Major components of this release:

Yahoo/Apache Hadoop 0.20-security-203 Apache Cassandra 0.8-beta2 Apache Hive 0.7 Apache Thrift 0.6.0 Monitoring & Managing the cluster (optional): DataStax OpsCenter v1.1
--

Please see respective components for their release notes specific to the version.

Known Issues

For the latest on these issues: <https://datastax.jira.com/browse/BRISK>

BRISK-169	HiveSQL drop table hangs (infinite loop of queries in DEBUG) after performing lots of alters to a table or after inserting data into a table with an index
BRISK-172	Tests/Jobs that read mapred-site.xml directly will fail to find JobTracker
BRISK-173	AMI has intermittent issue with starting up raid devices
BRISK-175	NullPointerException running ALTER TABLE foo TOUCH
BRISK-176	Cleaner build for source download that connects major components modified on disk and remotely

BRISK-177

Centos does not support sed -E so modify HIVE script to sed -r

Product Download & Documents Links

Product downloads, updates in form of AML, Source and Binary Archives, .deb are available from DataStax website.

archive	https://github.com/riptano/brisk/archives/brisk1
debian	debian.datastax.com
rpm	coming soon to: rpm.datastax.com

<http://www.datastax.com/products/brisk>

Sources & build directions are hosted on github - We encourage your participation.

<https://github.com/riptano/brisk>
<https://github.com/riptano/hive>
<https://github.com/riptano/brisk-hadoop-common>

Getting Started documents, faqs and wikis are hosted at the DataStax developer center:

<http://www.datastax.com/dev/>

Feedback:

Please join us for feedback and conversations on Brisk email lists:

<http://groups.google.com/group/brisk-users>
<http://groups.google.com/group/brisk-dev>
<http://groups.google.com/group/brisk-commit>
and on Internet Relay Chat (irc): **#datastax-brisk**