## About Brisk

Brisk is an open-source Hadoop and Hive distribution developed by DataStax that utilizes Apache Cassandra for its core services and storage. Brisk provides Hadoop MapReduce capabilities using CassandraFS, an HDFS-compatible storage layer inside Cassandra. By replacing HDFS with CassandraFS, users are able to leverage their current MapReduce jobs on Cassandra's peer-to-peer, fault-tolerant, and scalable architecture. Brisk is also able to support dual workloads, allowing you to use the same cluster of machines for both real-time applications and data analytics without having to move the data around between systems.

Brisk is now available via Apache license v2.0. The DataStax team welcomes your valued feedback.

### Release Contents

* *Brisk is comprised of the following components. For component-specific information, refer to their respective release notes and documentation.*

- Apache Hadoop 0.20.203.0 + (HADOOP-7172, HADOOP-5759, HADOOP-7255)
- Cassandra 0.8.0 + (CASSANDRA-2683)
- Apache Hive 0.7
- Apache Pig 0.8.3

## New Features in Brisk 1.0 Beta 2

The following new features have been added in this release:

| Feature | Description |
|---|---|
| BRISK-12 | Apache Pig Integration. See the DataStax Documentation for more information about using Pig in Brisk. |
| BRISK-89 | Job Tracker Failover. See the DataStax Documentation for more information about using the new `brisktool movejt` command. |
| BRISK-207 | New Snappy Compression Codec built on Google Snappy is now used internally for automatic CassandraFS block compression. |
| BRISK-180 | Automap Cassandra Column Families to Hive Tables in the Brisk Hive Metastore. |
| BRISK-152 | Add a second HDFS layer in CassandraFS for long-term data storage. This is needed because the blocks column family in CFS requires frequent |

| Feature | Description |
|---|---|
| | compactions - Hadoop uses it during MapReduce processing to store small files and temporary data. Compaction cleans this temporary data up after it is not needed anymore. Now there is the `cassandra:///` and `cassandra-archive:///` HDFS layers within CFS. The blocks column family in `cassandra-archive:///` will have compaction disabled to improve performance for static data stored in CFS. |

## Major Fixes in Brisk 1.0 Beta 2

Brisk 1.0 Beta 2 also incudes the following major fixes. For details on all fixes in Beta 2, see the Brisk Jira Project Web site:

| Issue | Description |
|---|---|
| BRISK-126 | Remove multiple slf4j warnings |
| BRISK-203 | Use `batchMutate` instead of `insert` in `HiveCassandraOutputFormat` |
| BRISK-219 | Cassandra super columns not mapping in Hive |
| BRISK-220 | Improve performance of `hadoop fs -ls` |
| CASSANDRA-2683 | Compaction issue causing secondary index corruption. |

## Open Issues

For a description of the open issues in Brisk, see the Brisk Jira Project Web site.